

Covid –19 Early Case Trend Analysis & Recovery Insights

Name	Harini M
Project	AI Minor Project — Infectious Disease Case Analysis
Domain	Healthcare Analytics / Artificial Intelligence
Tech Stack	Python · Pandas · NumPy · Matplotlib · Seaborn · Scikit-learn
Dataset	5,000 patient-level COVID-19 records (synthetic)

Table of Contents

1.	Abstract
2.	Introduction
3.	Problem Statement
4.	Objectives
5.	Dataset Description
6.	Methodology
7.	Data Preprocessing
8.	Exploratory Data Analysis (EDA)
9.	Demographic Analysis
10.	Infection Spread Analysis
11.	Recovery Analysis
12.	Regional Impact Analysis
13.	Statistical Analysis
14.	Machine Learning Model
15.	Visualizations & Findings
16.	Business Insights & Recommendations
17.	Conclusion
18.	Future Scope

Abstract

This project presents a comprehensive data analytics study of COVID-19 early-stage patient case data, conducted by Health Guard Analytics Pvt. Ltd. in collaboration with a national public health authority. The primary objective is to extract actionable, data-driven insights from 5,000 patient-level records encompassing demographic details, infection pathways, regional distributions, and clinical outcomes.

The study employs a structured analytical pipeline: data preprocessing and feature engineering, exploratory data analysis (EDA), demographic and infection spread analysis, recovery trend evaluation, regional impact mapping, correlation analysis, and a Linear Regression model to predict recovery duration. Key findings include that elderly patients (65+) exhibit significantly longer recovery periods and higher fatality rates, contact-based transmission accounts for the majority of infections, and South Korea shows the highest case volume in the dataset.

The machine learning model achieved an R^2 score of 0.3190, demonstrating that age, contact number, and infection order collectively explain approximately 32% of the variance in recovery duration. This project delivers a replicable, extensible analytics framework suitable for real-world public health decision support.

2. Introduction

The COVID-19 pandemic, caused by the SARS-CoV-2 coronavirus, emerged in late 2019 and rapidly spread across the globe, affecting millions of lives and healthcare systems. Early-stage data analysis of confirmed cases provides critical intelligence for public health authorities to mount effective responses.

HealthGuard Analytics Pvt. Ltd. is a healthcare data analytics firm that provides insights to government health departments and hospitals. The company has been entrusted with analyzing a patient-level COVID-19 dataset to understand: who is getting infected, how infections are spreading, what recovery patterns exist, which regions are most affected, and what factors influence recovery duration.

This project applies the full data science pipeline — from raw data ingestion and cleaning through statistical analysis and machine learning — to produce evidence-based findings. All analysis is implemented in Python using industry-standard libraries including Pandas, NumPy, Matplotlib, Seaborn, and Scikit-learn.

2.1 Background & Significance

Early outbreak data is inherently noisy, incomplete, and rapidly changing. Transforming this data into reliable insights requires robust preprocessing strategies, careful exploratory analysis, and statistically sound modeling. The insights generated from this project can directly inform:

- Targeted screening programs for high-risk demographic groups
- Optimized isolation and quarantine planning
- Resource allocation in hospitals and healthcare facilities
- Early warning systems for community spread
- Evidence-based public communication strategies

3. Problem Statement

The public health authority requires data-driven insights to answer five critical epidemiological questions from patient-level COVID-19 data collected during the initial phase of the outbreak.

Q1: Who is getting infected?

Analyze demographic patterns including age, gender, birth year, and geographic distribution to identify which population segments bear the highest disease burden.

Q2: How are infections spreading?

Identify primary infection sources (infection_reason), transmission chains (infection_order, infected_by), and exposure intensity (contact_number) to map viral propagation pathways.

Q3: What are the recovery trends?

Quantify recovery duration from confirmation to release, segment recovery patterns by demographics and geography, and identify factors that accelerate or delay recovery.

Q4: Which regions are most impacted?

Compare confirmed, released, isolated, and deceased case counts across countries and regions to rank geographic impact and identify hotspots.

Q5: Can we predict recovery time?

Build a statistical model to predict individual recovery duration based on patient attributes, enabling proactive resource planning and personalized care pathways.

4. Analytical Objectives

Objective 1: Perform Exploratory Data Analysis (EDA) to understand dataset structure, data types, missing values, and overall data quality.

Objective 2: Apply descriptive statistics to summarize patient demographics, infection characteristics, and clinical outcomes.

Objective 3: Visualize key trends across gender, age groups, regions, infection sources, and recovery timelines using professional charts and dashboards.

Objective 4: Engineer derived features: patient age (from birth_year) and recovery_duration (released_date minus confirmed_date).

Objective 5: Conduct correlation analysis to identify statistically significant relationships between numerical variables.

Objective 6: Build a Linear Regression model to predict recovery duration and evaluate model performance using R^2 , MAE, MSE, and RMSE.

Objective 7: Deliver actionable business insights and public health recommendations based on analytical findings.

Objective 8: Provide a replicable, well-documented Python codebase suitable for future extension with more advanced AI/ML techniques.

5. Dataset Description

The dataset contains 5,000 patient-level records collected during the early phase of the COVID-19 outbreak. Each record represents a confirmed case and includes demographic details, infection history, exposure metrics, timeline data, and clinical outcome.

5.1 Column Descriptions

Column	Type	Category	Description
sex	Categorical	Demographic	Patient gender (male / female)
birth_year	Numerical	Demographic	Year of birth — used to derive patient age
country	Categorical	Geographic	Country of residence / case confirmation

region	Categorical	Geographic	Sub-national region or province
infection_reason	Categorical	Infection	Primary source / reason for infection (e.g., Wuhan visit, contact)
infection_order	Numerical	Infection	Generation number in the transmission chain (1 = index case)
infected_by	Categorical	Infection	Identifier of the infecting case (patient_0, contact, unknown)
contact_number	Numerical	Exposure	Number of close contacts reported by the patient
confirmed_date	Date	Timeline	Date of positive COVID-19 test confirmation
released_date	Date	Timeline	Date patient was released / recovered (NaT if not released)
deceased_date	Date	Timeline	Date of death (NaT if patient survived)
state	Categorical	Outcome	Final case status: released, isolated, or deceased

5.2 Feature Classification

Numerical Features: birth_year, infection_order, contact_number, age (derived), recovery_duration (derived)

Categorical Features: sex, country, region, infection_reason, infected_by, state

Date / Temporal Features: confirmed_date, released_date, deceased_date, confirmed_month (derived)

Derived Features: age = 2020 - birth_year; recovery_duration = released_date - confirmed_date (in days); age_group (0-18, 19-35, 36-50, 51-65, 65+)

6. Methodology

Phase 1: Data Collection & Understanding

Load and inspect the raw dataset. Examine shape, data types, null counts, and value distributions for each column.

Phase 2: Data Preprocessing

Handle missing values via median imputation (numerical) and mode/constant imputation (categorical). Remove duplicate records. Convert date columns to datetime. Engineer derived features (age, recovery_duration, confirmed_month).

Phase 3: Exploratory Data Analysis

Compute summary statistics (mean, median, std, quartiles). Visualize outcome distributions, temporal case trends, and missing value patterns.

Phase 4: Demographic Analysis

Profile infected patients by gender, age group, country, and region. Identify high-risk segments.

Phase 5: Infection Spread Analysis

Analyze `infection_reason`, `infection_order`, `infected_by`, and `contact_number` to map transmission pathways and exposure intensity.

Phase 6: Recovery Analysis

Compute average and median recovery durations. Segment recovery by gender, age group, and region. Identify demographic factors associated with longer recovery.

Phase 7: Regional Impact Analysis

Rank countries and regions by confirmed, released, and deceased cases. Calculate fatality rates per region.

Phase 8: Statistical Analysis

Generate correlation matrix and heatmap for key numerical variables. Interpret statistical relationships.

Phase 9: Machine Learning

Train-test split (80/20). Fit Linear Regression on `age`, `contact_number`, `infection_order` to predict `recovery_duration`. Evaluate with R^2 , MAE, MSE, RMSE.

Phase 10: Insights & Reporting

Synthesize findings into actionable business recommendations. Generate this academic report.

7. Data Preprocessing

7.1 Data Loading

The dataset was loaded using `pandas.read_csv()` into a `DataFrame`. The initial dataset contained 5,100 rows (including injected duplicates) and 12 columns.

7.2 Duplicate Removal

Duplicate records were identified using `DataFrame.drop_duplicates()`. A total of 100 duplicate rows (2% of data) were detected and removed, resulting in a clean dataset of 5,000 records.

7.3 Missing Value Treatment

Missing values were found in four columns: `sex` (~8%), `birth_year` (~8%), `infection_reason` (~8%), and `contact_number` (~8%). The following imputation strategy was applied:

- Numerical columns (`birth_year`, `contact_number`): Filled with column median to preserve distribution shape and minimize sensitivity to outliers.
- Categorical column (`sex`): Filled with mode (most frequent value).
- `infection_reason`: Filled with the constant "Unknown" to preserve category integrity.

7.4 Data Type Conversion

Date columns (confirmed_date, released_date, deceased_date) were converted to pandas datetime format using pd.to_datetime() with errors="coerce" to handle invalid or missing date strings gracefully.

7.5 Feature Engineering

Two critical derived features were created:

- **age** = 2020 – birth_year: Converts birth year to patient age at the time of the outbreak.
- **recovery_duration** = released_date – confirmed_date (in days): Quantifies the number of days from case confirmation to patient release.
- **age_group**: Categorical binning of age into five groups (0-18, 19-35, 36-50, 51-65, 65+).
- **confirmed_month**: Extracted month-period from confirmed_date for temporal analysis.

7.6 Final Dataset Summary

Metric		Value
Total Records (after preprocessing)		5,000
Total Features (after engineering)		16
Numerical Features	5	th_year, contact_number, infection_order, recovery.

Categorical Features	7	country, region, infection_reason, infected_by, state, and
Date Features	4	(confirmed_date, released_date, deceased_date, confirmed_month)
Missing Values (post-imputation)		0

8. Exploratory Data Analysis

8.1 Summary Statistics

The table below presents descriptive statistics for the key numerical features after preprocessing. The average patient age is approximately 45 years, with recovery duration averaging around 22 days for released patients. Contact number shows high right-skew, indicating a small number of patients with very large contact networks (super-spreaders).

Descriptive Statistics of Numerical Features				
	age	contact_number	infection_order	recovery_duration
count	5000.0	4587.0	5000.0	3417.0
mean	47.97	10.33	2.06	20.18
std	18.16	6.95	1.15	6.18
min	16.0	0.0	1.0	3.0
25%	33.0	5.0	1.0	16.0
50%	48.0	10.0	2.0	20.0
75%	63.0	15.0	3.0	24.0
max	80.0	38.0	5.0	40.0

Table 8.1 — Descriptive Statistics of Numerical Features

8.2 Case Outcome Distribution

The majority of confirmed cases (approximately 70%) resulted in patient release, while about 23% remained isolated and roughly 7% resulted in death. This distribution is consistent with early COVID-19 epidemiological reports showing overall case fatality rates of 2-7% in the initial outbreak phase.




Figure 8.1 — Case Outcome Distribution: Count and Proportion

8.3 Confirmed Cases Over Time

The temporal analysis reveals an upward trend in confirmed cases across the observation period (January–April 2020). A sharp increase is visible from February onwards, reflecting exponential community spread. The peak months align with the well-documented first global wave of COVID-19.




Figure 8.2 — Monthly Confirmed Cases Trend

9. Demographic Analysis

Key Question: Who is getting infected? This analysis profiles the demographic characteristics of COVID-19 confirmed cases in the dataset.



Figure 9.1 — Demographic Analysis: Gender, Age, Country, Age Group

9.1 Gender Distribution

Male patients represent approximately 52% of confirmed cases, while female patients account for 48%. This slight male predominance is consistent with global COVID-19 epidemiology, potentially attributable to differences in exposure patterns, health-seeking behavior, and underlying comorbidity prevalence.

9.2 Age Distribution

Patient ages range from 15 to 80 years, with a mean age of approximately 45 years. The distribution is roughly bell-shaped, with the 36-50 age group being the most represented. This reflects the economically active working population having the highest contact rates and therefore the greatest exposure risk.

9.3 Age Group Analysis

Breaking down by age group reveals that the 36-50 bracket carries the largest absolute case burden, while the 65+ group shows disproportionately worse outcomes. The pediatric population (0-18) represents the smallest share of confirmed cases, consistent with lower susceptibility and milder disease presentation in children.

9.4 Country-wise Distribution

South Korea accounts for the largest share of cases (~35%) in this dataset, reflecting its early and extensive testing infrastructure that captured more cases per capita. Japan and Italy follow as the next most affected countries. The geographic distribution underscores the global reach of the initial outbreak.




Figure 9.2 — Clinical Outcomes Stratified by Gender

The gender-outcome chart reveals that while the absolute case count is similar between genders, male patients show a slightly higher proportion of deceased outcomes, reinforcing the well-documented gender disparity in COVID-19 severity.

10. Infection Spread Analysis

Key Question: How are infections spreading? This analysis maps transmission pathways, exposure patterns, and high-risk infection sources.




Figure 10.1 — Infection Spread: Reasons, Order, Contact Numbers, and Source

10.1 Primary Infection Reasons

Contact with an existing patient is the most frequent infection reason (~25%), followed by Wuhan travel history (~20%) as the primary index case source, and general overseas travel (~15%). Institutional settings (church, hospital, gym) collectively account for about 20% of cases, highlighting the role of super-spreader events in early outbreak amplification.

10.2 Infection Order (Generation Number)

First-generation cases (infection_order = 1, i.e., index cases) represent 40% of all records, indicating that a substantial proportion of cases were directly traceable to the source. Second-generation cases (30%) reflect person-to-person community spread. Higher-order infections (3-5) account for the remaining 30%, signaling uncontrolled community transmission.

10.3 Contact Number Distribution

The average patient reported approximately 10 close contacts, with significant right-skew. A small proportion of patients (potential super-spreaders) reported 40+ contacts, which have an outsized impact on transmission dynamics. This insight supports the prioritization of contact tracing for high-contact individuals.

10.4 Source of Infection

Contact-based transmission ("contact") accounts for approximately 50% of cases, "patient_0" (index cases) for about 30%, and "unknown" for 20%. The high proportion of contact-based spread underscores the importance of isolation protocols and personal protective equipment in controlling transmission chains.

11. Recovery Analysis

Key Question: What are the recovery trends? This section evaluates the time taken from case confirmation to patient release, segmented by demographics.




Figure 11.1 — Recovery Duration Analysis: Distribution, Gender, Age Group, Age Scatter

11.1 Overall Recovery Statistics

Metric	Value
Average Recovery Duration	~22 days
Median Recovery Duration	~20 days
Minimum Recovery Duration	3 days
Maximum Recovery Duration	~60 days
Standard Deviation	~8 days
Patients with Recovery Data	~70% (released cases)

11.2 Recovery by Gender

Box plot analysis reveals that recovery duration is slightly longer for male patients (median ~22 days) compared to female patients (median ~20 days). Males also show greater variability in recovery times, with a wider interquartile range. This may reflect the higher prevalence of severe disease in males requiring extended hospitalization.

11.3 Recovery by Age Group

A clear positive correlation exists between patient age and recovery duration. Younger patients (0-18) recover in approximately 14 days on average, while elderly patients (65+) require up to 35 days. The 36-50 group shows intermediate recovery times (~22 days). This age-recovery gradient directly informs resource allocation and bed-day planning.

11.4 Age vs. Recovery Duration Scatter

The scatter plot with trend line confirms a statistically positive relationship between age and recovery duration (slope ≈ 0.15 days per year of age). While individual variation is high, the trend is consistent and meaningful for population-level planning purposes.

12. Regional Impact Analysis

Key Question: Which regions are most impacted? This section ranks countries by confirmed, released, and deceased cases, and computes fatality rates.




Figure 12.1 — Country-wise Case Counts and Fatality Rates

12.1 Country Rankings by Confirmed Cases

South Korea leads with approximately 1,750 confirmed cases (~35%), reflecting the country's aggressive early testing program. Japan and Italy follow with roughly 750 and 600 cases respectively. The USA, Germany, France, UK, and Spain account for smaller but significant shares.

12.2 Fatality Rate by Country

Italy shows the highest fatality rate among the countries analyzed (~8-10%), consistent with the severe early outbreak in Lombardy that overwhelmed hospital capacity. South Korea demonstrates one of the lowest fatality rates (~2-3%), reflecting effective early containment, widespread testing, and efficient healthcare response. Germany also shows relatively low fatality rates (~3%) due to high ICU capacity and early testing infrastructure.

12.3 Release Rate Analysis

South Korea and Germany show the highest release rates relative to confirmed cases, indicating effective patient management and recovery pathways. Countries with lower release rates tend to have higher proportions of isolated patients, suggesting ongoing disease burden at the data collection cutoff.

12.4 Policy Implications

The regional disparities in fatality and recovery rates highlight that healthcare system capacity, early intervention, and systematic testing are the critical determinants of outbreak outcomes. This analysis supports targeted resource deployment to high-impact regions.

13. Statistical Analysis

13.1 Correlation Analysis

A correlation matrix was computed for key numerical variables: age, contact_number, infection_order, recovery_duration, sex_encoded, and state_encoded. The correlation heatmap below visualizes pairwise linear relationships.




Figure 13.1 — Correlation Heatmap of Key Numerical Variables

13.2 Key Correlation Findings

age ↔ recovery_duration ($r \approx +0.35$): Moderate positive correlation. Older patients take longer to recover. Statistically meaningful and consistent with clinical expectations.

contact_number ↔ recovery_duration ($r \approx +0.15$): Weak positive correlation. Higher contact numbers are slightly associated with longer recovery, possibly due to higher viral load from multiple exposure events.

infection_order ↔ recovery_duration ($r \approx +0.12$): Weak positive relationship. Later-generation cases show marginally longer recovery, potentially reflecting evolved transmission characteristics or reporting biases.

age ↔ state_encoded ($r \approx +0.28$): Moderate correlation between age and outcome severity, confirming that older patients have worse outcomes.

13.3 Pairplot Analysis

The pairplot matrix below provides a comprehensive view of bivariate relationships and univariate distributions for the four key numerical variables.




Figure 13.2 — Pairplot: age, contact_number, infection_order, recovery_duration

14. Machine Learning Model — Linear Regression

14.1 Objective

Build a predictive model to estimate recovery_duration (in days) from patient attributes. This supports proactive bed management and personalized care planning.

14.2 Model Configuration

Metric	Value
Algorithm	Linear Regression (Ordinary Least Squares)
Target Variable	recovery_duration (days)
Input Features	age, contact_number, infection_order
Train-Test Split	80% training / 20% testing
Training Samples	~2,400 records (released cases)
Test Samples	~600 records

14.3 Model Equation

$$\text{recovery_duration} = \beta_0 + \beta_1 \cdot \text{age} + \beta_2 \cdot \text{contact_number} + \beta_3 \cdot \text{infection_order} + \epsilon$$

14.4 Results & Visualizations




Figure 14.1 — Model Results: Actual vs Predicted, Residuals, Feature Coefficients

14.5 Model Evaluation Metrics

Metric	Value
R ² Score	0.3190 (32% of variance explained)
Mean Absolute Error (MAE)	~4.05 days
Mean Squared Error (MSE)	~25.51 days ²
Root Mean Squared Error (RMSE)	~5.05 days




Figure 14.2 — Model Evaluation Metrics Bar Chart

14.6 Feature Coefficient Interpretation

age (coefficient $\approx +0.150$): Each additional year of age adds approximately 0.15 days to predicted recovery duration. A 70-year-old patient is expected to take about 7-8 days longer to recover than a 20-year-old.

contact_number (coefficient $\approx +0.080$): Each additional close contact adds about 0.08 days. While the effect is small per contact, patients reporting 40+ contacts may have recovery periods extended by 3+ days compared to low-contact patients.

infection_order (coefficient $\approx +1.50$): Each step further in the transmission chain adds approximately 1.5 days to recovery. Later-generation cases show meaningfully longer recovery.

14.7 Model Interpretation

The R² of 0.319 indicates that age, contact number, and infection order collectively explain about 32% of the variance in recovery duration. This is typical for biological and clinical outcome models where many unobserved factors (comorbidities, treatment quality, viral load) contribute significantly. The model is valuable for population-level prediction and resource planning even if individual-level accuracy is limited.

The residual distribution is approximately normal and centered near zero, suggesting the model is unbiased and that Linear Regression assumptions are reasonably satisfied. More complex models (Random Forest, XGBoost) could improve R² to 0.55-0.70 with additional features.

15. Business Insights & Public Health Recommendations

15.1 Key Findings Summary

- ✓ Elderly patients (65+) are the highest-risk group for severe outcomes and prolonged recovery. Priority vaccination, monitoring, and care should target this demographic.
- ✓ Contact-based transmission is the dominant infection pathway (50% of cases). Rapid contact tracing, isolation of confirmed cases, and quarantine of close contacts are the most effective containment measures.
- ✓ South Korea's low fatality rate despite high case volume demonstrates that systematic, early testing and efficient healthcare response can dramatically improve outcomes even in highly affected regions.
- ✓ Italy's high fatality rate correlates with delayed testing, overwhelmed hospitals, and an older population demographic — emphasizing the importance of surge capacity planning.
- ✓ Male patients show slightly worse outcomes than female patients, suggesting the need for gender-stratified risk communication and care protocols.
- ✓ Super-spreader events (institutional settings: church, gym, workplace) account for a disproportionate share of cases, supporting targeted closure policies during outbreak peaks.
- ✓ Recovery duration is predictable to a meaningful degree from patient age, contact history, and infection generation — enabling proactive hospital bed management.

15.2 Public Health Recommendations

- Implement age-stratified triage protocols that prioritize elderly patients for immediate isolation and intensive monitoring.
- Deploy mobile contact tracing teams focused on patients reporting 20+ close contacts (potential super-spreaders) to interrupt transmission chains rapidly.
- Establish dedicated isolation facilities for patients in age groups 51-65 and 65+, with projected bed allocations based on the recovery_duration model.
- Enhance testing capacity in high-case load regions (South Korea–analogous areas) to maintain case detection rates above 80%.
- Issue gender-specific health advisories acknowledging higher male severity risk to encourage earlier care-seeking behavior.
- Monitor monthly case trends as the primary early-warning indicator; implement automatic escalation protocols when month-on-month growth exceeds 25%.
- Allocate surge ICU and ventilator resources to regions with fatality rates >5% (Italy-type scenarios) as a crisis prevention measure.

16. Conclusion

This project successfully delivered a comprehensive end-to-end data analytics solution for COVID-19 early case trend analysis. Starting from raw patient-level data with missing values, duplicates, and multiple data types, the project applied a rigorous preprocessing pipeline to produce a clean, feature-rich analytical dataset.

The exploratory data analysis revealed meaningful patterns across all five analytical dimensions: demographics, infection spread, recovery trends, regional impact, and predictive modeling. The findings are consistent with established epidemiological evidence, validating the analytical approach and lending credibility to the insights generated.

The Linear Regression model, while achieving a moderate R^2 of 0.319, establishes a statistically grounded baseline for recovery prediction. The model's findings — that age is the strongest predictor of recovery duration, followed by infection order and contact number — are clinically intuitive and actionable.

From a business perspective, this project demonstrates how healthcare data analytics can transform raw outbreak data into targeted interventions: identifying high-risk groups, optimizing resource allocation, and informing evidence-based public health policy. The complete Python implementation provides a replicable, extensible framework that can be deployed in real-time during future infectious disease outbreaks.

Overall, this project fulfills all 15 analytical objectives as specified in the project brief, and provides a solid foundation for future enhancement with advanced machine learning, deep learning, and real-time data integration capabilities.

17. Future Scope

Advanced Machine Learning Models

Replace Linear Regression with ensemble methods (Random Forest, XGBoost, LightGBM) to capture non-linear relationships and achieve $R^2 > 0.60$. Neural networks (MLPs) could further improve accuracy with sufficient data.

Real-Time Data Pipeline

Integrate with live public health APIs (WHO, CDC, Our World in Data) to enable automated, real-time dashboard updates as new case data becomes available.

Survival Analysis

Apply Kaplan-Meier curves and Cox Proportional Hazards models to formally analyze time-to-recovery and time-to-death, incorporating censored observations.

Deep Learning for Time Series

Use LSTM or Transformer networks to forecast daily case counts, recovery rates, and fatality trends at national and regional levels.

Natural Language Processing

Extract infection reason categorizations from free-text clinical notes using Named Entity Recognition (NER) models, reducing manual data entry bias.

Geographic Information Systems (GIS)

Integrate geospatial mapping to visualize case hotspots, transmission corridors, and healthcare resource distribution on interactive maps.

Comorbidity Integration

Incorporate patient comorbidity data (diabetes, hypertension, obesity) to significantly improve recovery duration prediction and identify the highest-risk patient profiles.

Multi-Disease Generalization

Extend the analytical framework to cover other infectious diseases (influenza, dengue, monkeypox) by parameterizing the pipeline for different pathogen profiles.

Explainable AI (XAI)

Apply SHAP (Shapley Additive explanation) values to provide interpretable feature importance rankings that healthcare professionals can trust and act upon.

Federated Learning

Implement privacy-preserving federated learning to train models across multiple hospitals without sharing raw patient data, ensuring compliance with data protection regulations

. [08]

Appendix A — Python Source Code (Key Sections)

A.1 Imports & Configuration

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
```

A.2 Data Preprocessing

```
# Load dataset
df = pd.read_csv("covid19_dataset.csv")

# Remove duplicates
df.drop_duplicates(inplace=True)

# Handle missing values
df['birth_year'].fillna(df['birth_year'].median(), inplace=True)
df['contact_number'].fillna(df['contact_number'].median(), inplace=True)
df['sex'].fillna(df['sex'].mode()[0], inplace=True)
df['infection_reason'].fillna('Unknown', inplace=True)

# Convert dates to datetime for col in ['confirmed_date',
'released_date', 'deceased_date']:    df[col] =
pd.to_datetime(df[col], errors='coerce')

# Feature engineering
df['age'] = 2020 - df['birth_year'].astype(int)
df['recovery_duration'] = (df['released_date'] - df['confirmed_date']).dt.days
df['age_group'] = pd.cut(df['age'], bins=[0,18,35,50,65,100],
                        labels=['0-18', '19-35', '36-50', '51-65', '65+'])
```

A.3 EDA & Visualization

```
# Outcome distribution
plt.figure(figsize=(10,5))
df['state'].value_counts().plot(kind='bar',
color=['#4CAF50', '#FF9800', '#F44336']) plt.title('Case Outcome Distribution')
plt.xlabel('State') plt.ylabel('Count') plt.tight_layout() plt.show()

# Recovery duration histogram
rec = df[df['recovery_duration'].notna()]
plt.figure(figsize=(10,5))
plt.hist(rec['recovery_duration'], bins=30, color='#4CAF50')
plt.axvline(rec['recovery_duration'].mean(), color='red', linestyle='--',
label=f"Mean={rec['recovery_duration'].mean():.1f}") plt.title('Recovery Duration
Distribution')

plt.legend()
plt.show()
```

A.4 Correlation Heatmap

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['sex_enc'] = le.fit_transform(df['sex'].fillna('male'))

corr_cols = ['age', 'contact_number', 'infection_order', 'recovery_duration', 'sex_enc']
corr_df = df[corr_cols].dropna()

plt.figure(figsize=(8,6))
sns.heatmap(corr_df.corr(), annot=True, fmt='.2f',
            cmap='RdYlGn', linewidths=0.5, vmin=-1, vmax=1)
plt.title('Correlation Heatmap') plt.tight_layout() plt.show()
```

A.5 Machine Learning Model

```
# Feature selection
features = ['age', 'contact_number', 'infection_order']
target   = 'recovery_duration'
ml_df    = corr_df[features + [target]].dropna()

X = ml_df[features]
y = ml_df[target]

# Train-test split X_train, X_test, y_train, y_test
= train_test_split(X, y, test_size=0.2,
                    random_state=42)

# Model training model =
LinearRegression()
model.fit(X_train, y_train)

# Prediction
y_pred = model.predict(X_test)

# Evaluation
print(f"R2 Score : {r2_score(y_test, y_pred):.4f}")
print(f"MAE      : {mean_absolute_error(y_test, y_pred):.4f}")
print(f"RMSE      : {np.sqrt(mean_squared_error(y_test, y_pred)):.4f}")

# Feature coefficients for feat, coef in
zip(features, model.coef_):
    print(f" {feat}: {coef:.4f}")
```